Budget-Aware Self-Configuring Neural Architecture Search for Unseen Data

Sebastian Seifert, Luis Melzer & Leander von Gottberg

1 Motivation

The NAS Unseen-Data Challenge (AutoML 2026) evaluates a fully automated pipeline on *hidden* classification datasets under a strict wall-clock budget per dataset. Inputs arrive as image-shaped tensors but range from natural and satellite images to *encoded* data (rendered text, Sudoku boards, chess positions) where the spatial prior of a CNN misleads. All decisions must be made at runtime, blind.

Our starting point is the winning method of the 2025 edition, EGNAS [**?**], a macro–micro search (depth, width, per-layer operation) over convolutional blocks. Reproducing it surfaced two systematic weaknesses. First, rankings from short training runs are unreliable: operations that lead after a few epochs often lose after full training, so fine-grained (micro) operation search tends to *harm* final accuracy. Second, the space is confined to one architecture family: in our preliminary experiments a *fixed* 2-D dilated-convolution backbone in the style of temporal convolutional networks [**?**] clearly outperformed the fully searched CNN baseline on most datasets, while plain CNNs and randomly sampled architectures won precisely on the encoded datasets, in line with random search being a strong NAS baseline [**?**, **?**]. The leverage lies in choosing the right architecture *family* per dataset, not in fine-grained operation search.

2 Research Questions

RQ1. Can macro–micro NAS [**?**] be generalised across heterogeneous architecture families (plain CNNs, 2-D dilated and 1-D temporal convolution networks), making families *comparable* within a single time-budgeted search stage despite the known unreliability of low-fidelity rankings [**?**]? **RQ2.** Can cheap, training-free measurements of the raw data decide at runtime which architectural prior fits an unseen dataset? **RQ3.** How can harmful data augmentation be detected at runtime, given that augmentation acts as regularisation (it costs accuracy early and pays late) and that, as we observe, its decisions do not transfer between architecture families?

3 Approach

We propose a self-configuring pipeline that measures the data first, gates every decision, and spends all remaining time on training (Fig. **??**, p. 2). (1) *Measure*: a fingerprint computed directly from the raw input tensors, with no training involved. The mean correlation of neighbouring pixels separates the two data families with a wide margin: between -0.05 and 0.07 on the four encoded public datasets versus 0.38 to 1.00 on all others. (2) *Macro stage*: for image-like data the dilated backbone is selected directly. Only encoded data pays for a brief bake-off of all families at matched wall-clock. The leader is adopted only if it beats the robust default by more than the measured ranking noise, otherwise the pipeline *abstains* to the default, which makes every decision bounded-downside. (3) *Augmentation probe*: a two-arm A/B test run *after* the family choice with arms of the chosen family. It deviates from the default only on a clear margin (RQ3). (4) *Micro stage*: depth search anchored on an incumbent default. A deviation must win significantly *and* survive verification by re-training at a longer horizon. (5) A *random-search arm* over a wide operation vocabulary runs only on encoded data and only if budget remains [**?**]. (6) *Train to budget*: the winner trains on the full remaining time with a cosine learning-rate schedule [**?**] sized to reach zero exactly when the budget ends. All stage budgets are fractions of the clock, so the pipeline degrades gracefully from many hours down to minutes.

4 Required Steps

(1) Reproduce the EGNAS baseline and probe candidate architecture families at full budget (done). (2) Validate each decision mechanism in isolation against known full-budget outcomes (gate reliability, fingerprint separation, augmentation transfer). (3) Integrate the pipeline and verify all execution paths, including the low-budget collapse behaviour. (4) Full multi-seed evaluation. (5) A budget-ladder study (0.5h to 8h) measuring which decision stages pay at which budget. (6) A generalisation test on untouched external datasets (CIFAR-10/100, EMNIST, MedMNIST, spectrograms). (7) Packaging and stress tests.

5 Evaluation Strategy

Primary evaluation uses the 13 public competition datasets with the official adjusted score (accuracy versus a per-dataset benchmark, capped), at least three seeds for final configurations. Baselines: the reproduced EGNAS search, each fixed family, and a best-family oracle bounding the achievable routing gain. Every stage writes its measurements, margins, and choices to a per-dataset decision record, so gains and losses are attributable to individual decisions. Robustness is evaluated across budgets (30 minutes to 12 hours), on the external suite, and against published competition bests [**?**]. First results of the current pipeline, including the configuration it selects per dataset, are shown in Fig. **??** (p. 3).

6 Potential Extensions

The pipeline's ambition can scale with the budget: additional families (e.g. recurrent networks, which need longer probes to rank fairly) join the bake-off only at large budgets, and when families tie statistically, the surplus trains a heterogeneous ensemble, kept only if it beats the best single model. In parallel we create competition-style datasets blind to the pipeline (QR codes with partly covered regions, encoded board games) as a further generalisation test.

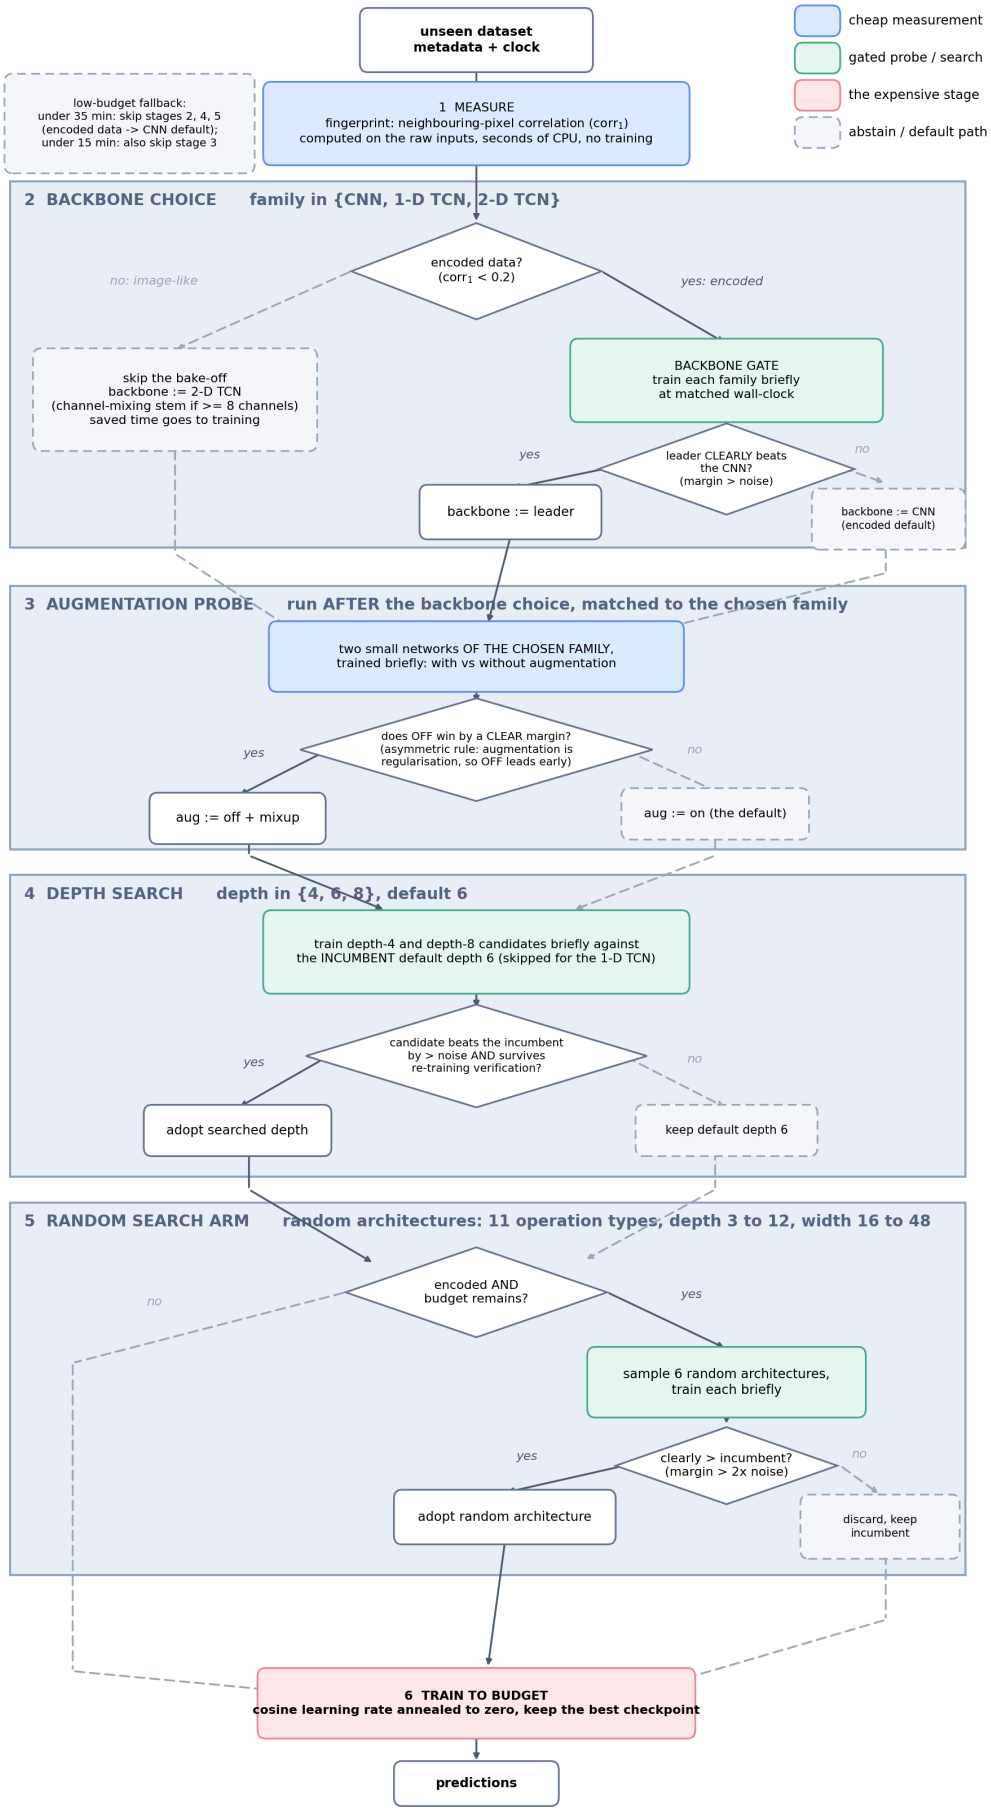


Figure 1: The decision pipeline, run once per unseen dataset. Diamonds are runtime decisions, dashed paths are the abstain-to-default routes.

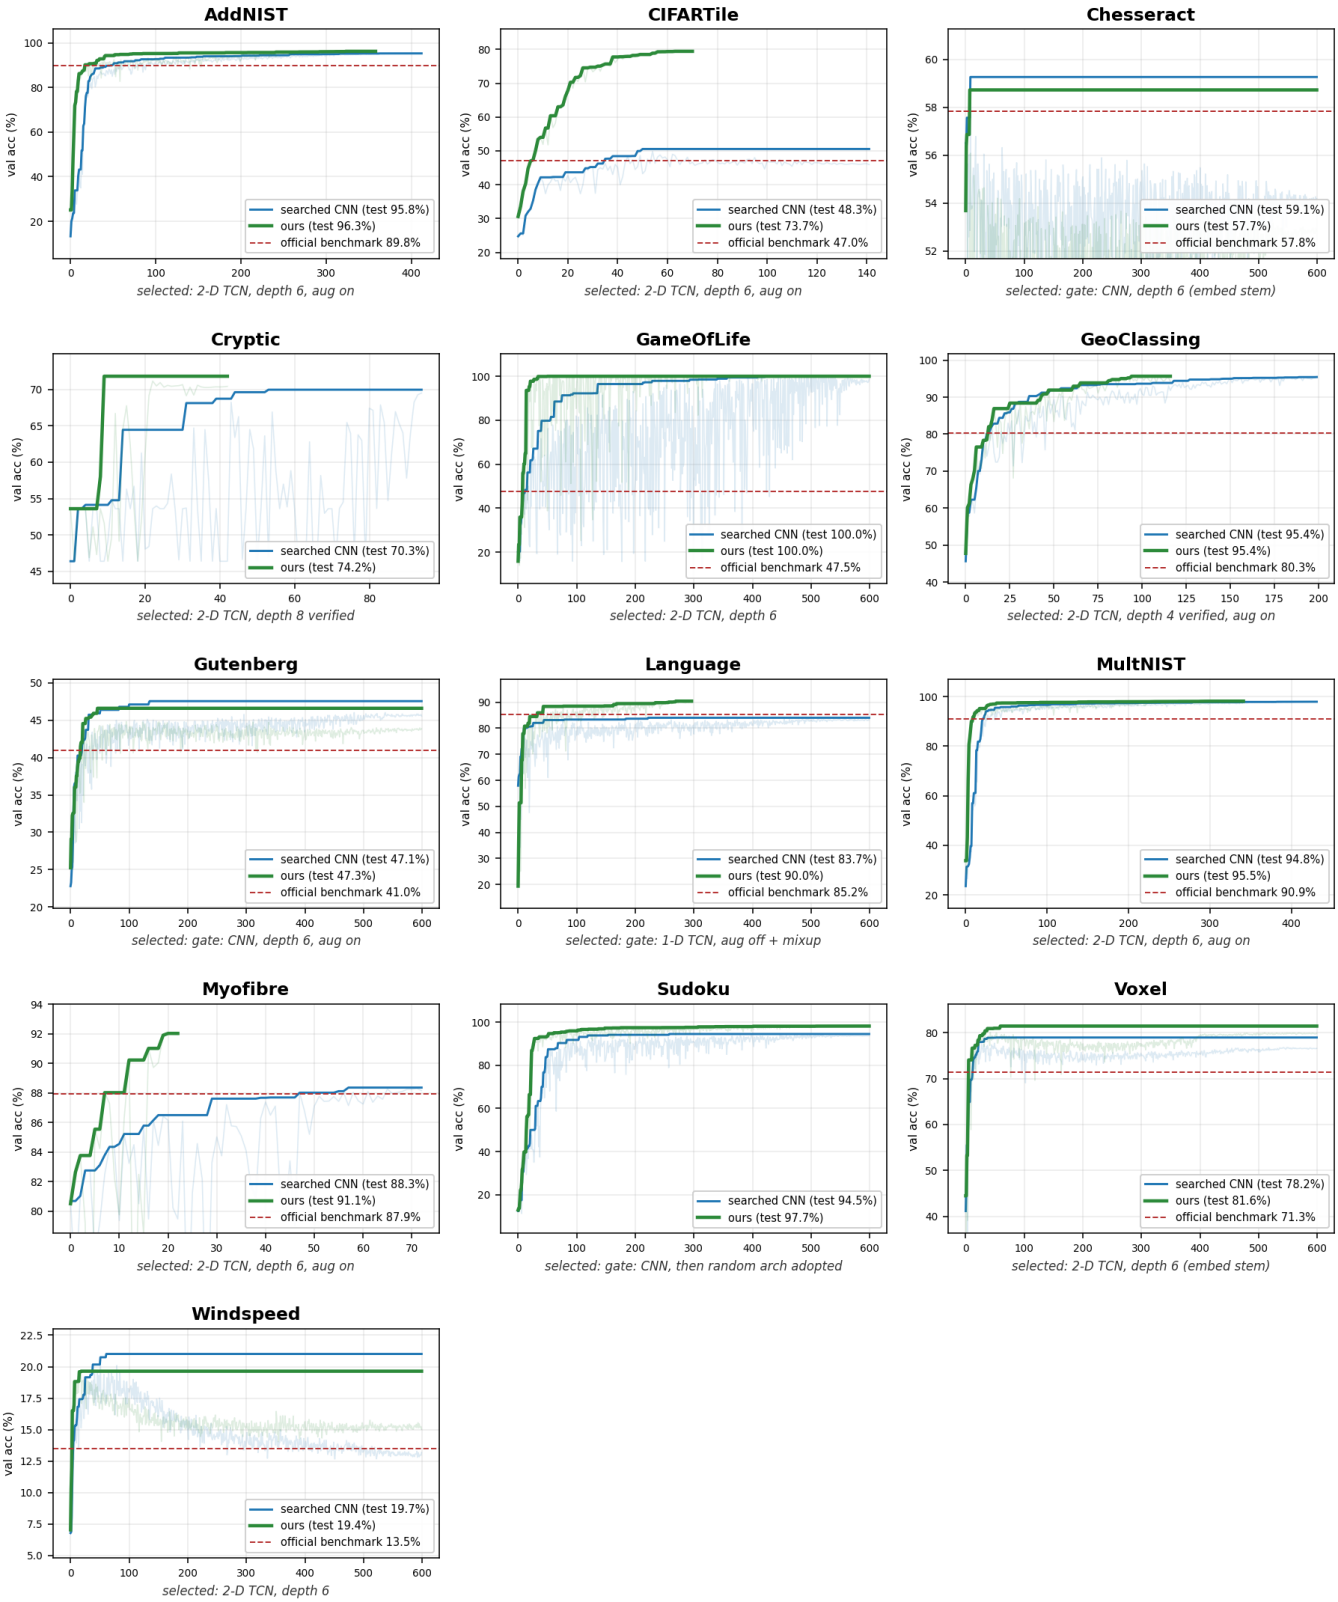


Figure 2: Our current best approach versus the searched CNN (our reproduction of the winning macro–micro NAS [**?**] with identical training improvements), 13 public datasets, single seed, 2 h each. Curves: validation accuracy (faint per-epoch, solid best-so-far). Legends: final test accuracy, which can differ from validation. Red dashed: the official benchmark from the dataset metadata; unpublished for Sudoku and Cryptic. The selected configuration is noted under each panel. Totals (adjusted score): ours **+62.5**, searched CNN +48.2.

References

[1] S. Siddiqui, C. Kyrkou, T. Theocharides. *Efficient Global Neural Architecture Search*. SN Computer Science 6:282, 2025.

[2] R. Geada, D. Towers, M. Forshaw, A. Atapour-Abarghouei, A. S. McGough. *Insights from the Use of Previously Unseen Neural Architecture Search Datasets*. CVPR 2024.

[3] S. Bai, J. Z. Kolter, V. Koltun. *An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling*. arXiv:1803.01271, 2018.

[4] L. Li, A. Talwalkar. *Random Search and Reproducibility for Neural Architecture Search*. UAI 2019.

[5] C. White, A. Zela, B. Ru, Y. Liu, F. Hutter. *How Powerful are Performance Predictors in Neural Architecture Search?* NeurIPS 2021.

[6] I. Loshchilov, F. Hutter. *SGDR: Stochastic Gradient Descent with Warm Restarts*. ICLR 2017.